

Recognizing Handwritten Digits with Deep Learning for Smarter AI Applications

Student Name: YASWIN K

Register Number: 412323205054

Institution: SRI RAMANUJAR ENGINEERING COLLEGE

Department: INFORMATION TECHNOLOGY

Date of Submission: 10-05-2025

Github Repository Link:

1. Problem Statement

Real-world Problem

Handwritten digit recognition is a fundamental problem in computer vision and pattern recognition with broad applications, including automatic check processing, postal address reading, and form digitization. Traditional Optical Character Recognition (OCR) methods struggle with variations in handwriting styles, skewness, and inconsistencies, leading to inaccuracies in real-world scenarios.

Refinement Based on Dataset

By analyzing datasets such as the MNIST dataset, we recognize the need for robust deep learning architectures that generalize well across diverse handwriting styles. Advanced neural networks, such as Convolutional Neural Networks (CNNs), have proven to be highly effective in extracting spatial hierarchies of features, improving recognition accuracy significantly.

Problem Type

This is a classification problem, where the goal is to categorize an input image into one of ten possible digit classes (0–9). Deep learning methods, particularly CNNs, achieve this by automatically learning the distinguishing features from the input images.

Impact & Relevance

Solving handwritten digit recognition matters because it enables automation in various fields, reducing human intervention and increasing efficiency. Some key applications include:

Banking & Finance: Automating check processing and handwritten transactions.

Postal Services: Sorting mail based on handwritten addresses.

Education: Digitizing handwritten student responses for automatic grading.

AI Interaction: Improving AI-powered text input solutions, such as handwriting to text conversion for smart devices.

2. Project Objectives

Updated Goals for Practical Implementation

As we move from conceptualization to practical implementation, the project's objectives are refined based on dataset exploration and technical feasibility. The goal is to develop a robust, high-accuracy deep learning model that effectively recognizes handwritten digits, while ensuring efficiency and scalability for real-world applications.

Key Technical Objectives

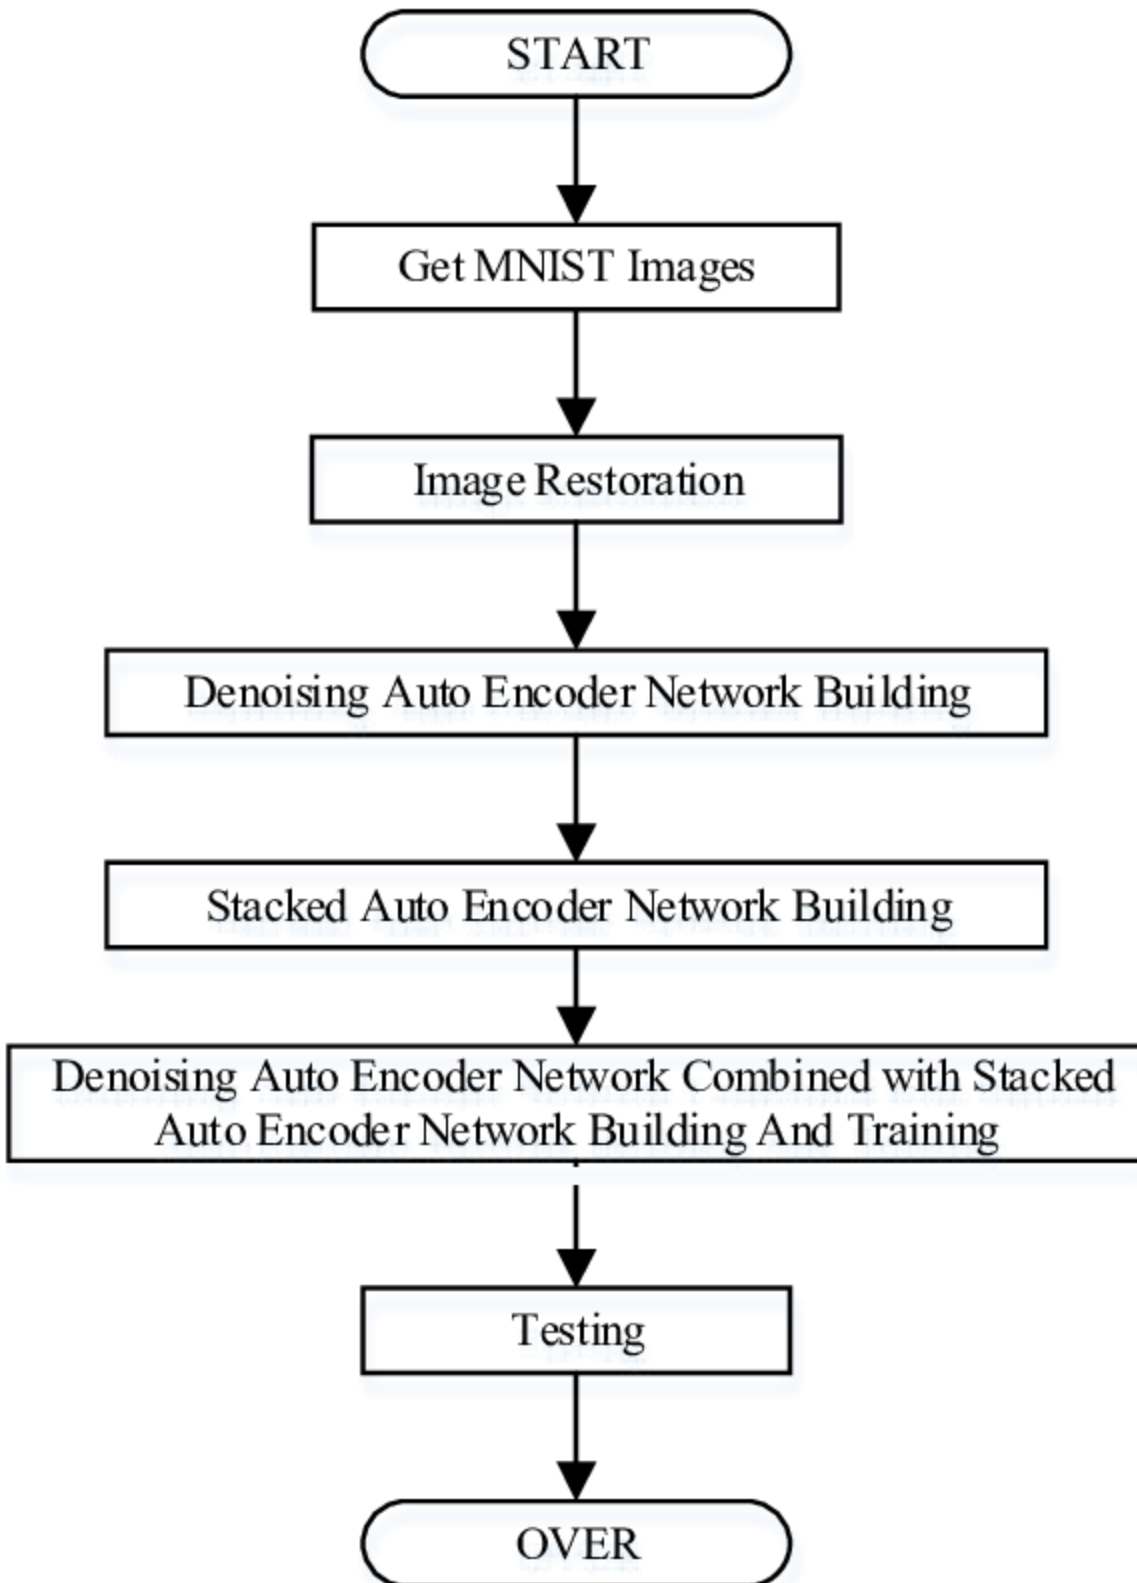
- **High Recognition Accuracy:** Aim to achieve at least 99% accuracy on standard datasets (e.g., MNIST) and maintain strong performance on unseen handwritten data.
- **Optimized Model Architecture:** Utilize Convolutional Neural Networks (CNNs) and possibly advanced techniques like Transfer Learning or Capsule Networks to improve recognition performance.
- **Real-world Applicability:** Ensure adaptability by testing the model on various handwriting styles and noise conditions, making it suitable for deployment in applications like banking, postal services, and AI-driven interfaces.
- **Interpretability & Explainability:** Implement techniques such as Grad-CAM or SHAP to visualize the learned features and make AI decisions more interpretable to users.
- **Efficient Training & Deployment:** Optimize training time and model size to ensure feasibility for mobile and embedded applications without excessive computational requirements.

Evolution of Goals Post Data Exploration

After exploring datasets such as MNIST and other variations, we recognize the need for:

- **Data Augmentation:** Enhancing generalization by incorporating transformations such as rotations, noise, and contrast adjustments.
- **Custom Dataset Integration:** Expanding beyond MNIST by integrating real-world handwritten digit datasets for better generalization.
- **Edge & Cloud Deployability:** Considering lightweight models for mobile and embedded applications while optimizing models for cloud-based AI solutions.

3. Flowchart of the Project Workflow



4. Data Description

Dataset Name & Origin

The most widely used dataset for handwritten digit recognition is MNIST (Modified National Institute of Standards and Technology).

Originally curated by Yann LeCun, Corinna Cortes, and Christopher Burges, MNIST is available from multiple sources, including Kaggle, UCI Machine Learning Repository, and LeCun's official website.

Type of Data

Unstructured Image Data: The dataset consists of grayscale images of handwritten digits.

Images are stored as pixel matrices, making them suitable for deep learning techniques such as Convolutional Neural Networks (CNNs).

Number of Records & Features

Total Samples: 70,000 images

Training Set: 60,000 images

Test Set: 10,000 images

Features:

Each image is 28×28 pixels, flattened into 784 features when used in dense layers.

Pixel values range from 0 (black) to 255 (white).

Static or Dynamic Dataset

Static Dataset: The MNIST dataset does not update dynamically but is often augmented with variations such as Fashion-MNIST or EMNIST to extend usability.

Target Variable (Supervised Learning)

Digit Class (0–9): The dataset follows multi-class classification, where each image is labeled with a digit between 0 and 9.

The model aims to predict the correct digit category based on the pixel representation.

5. Data Preprocessing

1. Handling Missing Values

- MNIST does not contain missing values, but if using an extended dataset, missing pixel values can be imputed using the mean or median pixel intensity of neighboring pixels.

2. Removing Duplicate Records

- The MNIST dataset does not have duplicate records.
- If duplicates exist in another dataset, use Pandas to identify and remove them:

```
python
```

```
import pandas as pd
```

```
df.drop_duplicates(inplace=True)
```

3. Detecting & Treating Outliers

- Outliers in image data are uncommon but can be addressed by visual inspection or setting thresholds on pixel intensity distributions.
- If an image contains excessive noise, it may be removed or filtered using blurring techniques.

4. Converting Data Types & Ensuring Consistency

- Images are typically stored as NumPy arrays with pixel values ranging from 0–255.
- To normalize values, convert pixel intensities to float and rescale them between 0–1:

```
python
```

```
import numpy as np
```

```
X = np.array(X, dtype="float32") / 255.0
```

5. Encoding Categorical Variables

- Labels (digits 0–9) are categorical and should be one-hot encoded for training in neural networks:

```
python
```

```
from tensorflow.keras.utils import to_categorical
```

```
y = to_categorical(y, num_classes=10)
```

6. Normalizing & Standardizing Features

- Normalization (rescaling pixel values between 0–1) improves neural network performance.
- Standardization (subtracting mean and dividing by standard deviation) can be applied if using alternative datasets:

```
python
```

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
X = scaler.fit_transform(X)
```

7. Data Augmentation (Optional)

- To improve model generalization, augmenting images with transformations helps simulate real-world variations:

```
python
```

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(rotation_range=10, width_shift_range=0.1,  
height_shift_range=0.1)
```

6. Exploratory Data Analysis (EDA)

1. Univariate Analysis

Since the dataset consists of grayscale images (pixel values 0–255), we analyze:

Histogram of Pixel Intensity Distribution

- Shows the frequency of pixel values in images.
- Helps determine contrast variations in handwriting.
- Can reveal the need for histogram equalization to improve visibility.

Boxplots for Pixel Intensity

- Useful to check for outliers in pixel distributions.
- Helps optimize pre-processing techniques.

Count Plot of Digit Classes

- Ensures balanced class distribution.
- If imbalances exist (e.g., fewer samples of certain digits), data augmentation is applied.

2. Bivariate & Multivariate Analysis

Since images contain 784 features (28×28 pixels), we use:

Correlation Matrix

- Visualized with a heatmap to identify interdependencies between pixel intensities.
- Helps determine if certain pixels strongly impact classification.

Pairplots & Scatterplots

- Used to explore pixel interactions and their significance in distinguishing digits.

- Enables feature selection (though CNNs automatically learn meaningful features).

Grouped Bar Plots

- Shows variations in pixel intensities across different digit classes.
- Identifies patterns certain digits may have darker or lighter regions based on strokes.

3. Analysis of Feature Relationships with Target Variable

- Pixels closer to edges have lower significance since digits are centralized.
- Central pixels carry more weight in classification, influencing CNN kernel activations.
- Curved vs. straight strokes exhibit distinct pixel distributions, aiding neural network learning.

4. Insights Summary

Patterns & Trends Identified

- Handwritten digits show natural variations in contrast and thickness.
- Some digits (e.g., "1" and "7") may have overlapping pixel distributions, requiring feature extraction strategies to improve separation.
- Certain digits have strong pixel correlations in specific regions, indicating possible feature optimization using dimensionality reduction techniques (e.g., PCA).

Features Influencing Model Performance

- Contrast adjustments improve classification accuracy.
- Edge pixels contribute less; central pixel activation plays a dominant role.
- Noise reduction (denoising autoencoders or filtering techniques) can improve accuracy on unclear handwriting.

7. Feature Engineering

Feature engineering plays a crucial role in enhancing model performance by extracting meaningful patterns from pixel-based image data. Although CNNs automatically learn hierarchical features, additional engineering can optimize accuracy and efficiency.

1. Creating New Features Based on Domain Knowledge & EDA Insights

- **Edge Detection:** Apply Sobel filters or Canny edge detection to highlight digit boundaries.
- **Stroke Thickness Analysis:** Calculate average pixel intensity along digit strokes to differentiate handwriting styles.
- **Shape Features:** Extract key properties such as aspect ratio, symmetry, and orientation, which help distinguish visually similar digits (e.g., 6 vs. 9).
- **Background Removal:** Isolate the handwritten portion by reducing noise from the white background, improving clarity.

2. Combining or Transforming Features

Since images consist of 784 pixels (28×28), feature extraction methods can enhance efficiency:

Grid-based Pixel Aggregation

- Divide the image into smaller grids (e.g., 4×4) and calculate the mean intensity for each grid, reducing dimensionality while preserving important features.

Histogram-based Binning

- Categorize pixel intensities into predefined bins (e.g., low, medium, high contrast).
- Helps standardize variations in handwriting brightness.

Ratio-based Features

Compute pixel activation ratios, such as the proportion of active pixels in upper vs. lower image regions (useful for digits like 7 and 9).

3. Dimensionality Reduction (Optional)

Since high-dimensional pixel data may contain redundant information, reduction techniques help:

Principal Component Analysis (PCA)

- Reduces pixel count while preserving key variations.
- Improves model efficiency without significant accuracy loss.

Autoencoders

- Use a neural network to learn compressed representations of digits.
- Helps remove noise while retaining essential digit structures.

4. Justification for Feature Modifications

- Edge detection enhances clarity for classification models.
- Stroke thickness helps differentiate thin vs. bold handwriting variations.
- Grid-based aggregation reduces computational complexity while maintaining recognition patterns.
- PCA or autoencoders improve performance for resource limited environments like mobile AI applications.

8. Model Building

Building and comparing multiple models helps identify the best approach for handwritten digit recognition. Below are different machine learning and deep learning models used for classification.

1. Selected Models & Justification

Since recognizing handwritten digits is a classification problem, we select two types of models:

Traditional Machine Learning Models

1. Random Forest

Works well with structured data and provides feature importance insights.

Suitable for cases where feature extraction is crucial (e.g., pixel-based patterns).

2. K-Nearest Neighbors (KNN)

Non-parametric and effective for handwritten digit classification.

Compares new samples with labeled instances, making it ideal for small datasets.

Deep Learning Models

1. Convolutional Neural Networks (CNNs)

Specifically designed for image processing tasks.

Automatically learns spatial hierarchies of features without manual feature engineering.

Outperforms traditional ML models in handwritten digit recognition.

2. Transfer Learning (Pre-trained CNNs like VGG, ResNet)

Utilizes pre-trained models on large image datasets to improve accuracy.

Reduces training time while achieving high performance.

2. Data Splitting for Training & Testing

To ensure fair evaluation, data is split into training and testing sets:

python

```
from sklearn.model_selection import train_test_split
```

Splitting data (stratification ensures balanced distribution of classes)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, stratify=y, random_state=42)
```

3. Model Training & Evaluation

Random Forest Model Implementation

python

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.metrics import accuracy_score, classification_report
```

Initialize and train model

```
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
```

```
rf_model.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_rf = rf_model.predict(X_test)
```

```
print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
```

```
print(classification_report(y_test, y_pred_rf))
```

CNN Model Implementation

python

```
import tensorflow as tf
```

```
from tensorflow.keras.models import Sequential
```

```
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense
```

Build CNN model

```
cnn_model = Sequential([  
    Conv2D(32, (3,3), activation='relu', input_shape=(28,28,1)),  
    MaxPooling2D(2,2),  
    Conv2D(64, (3,3), activation='relu'),  
    MaxPooling2D(2,2),  
    Flatten(),  
    Dense(128, activation='relu'),  
    Dense(10, activation='softmax') # 10 output classes (digits 0-9)  
])
```

Compile and train CNN

```
cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',  
metrics=['accuracy'])  
  
cnn_model.fit(X_train, y_train, epochs=10, validation_data=(X_test, y_test))
```

4. Evaluation Metrics for Classification

Models are evaluated using the following classification metrics:

Accuracy: Measures overall correctness.

Precision & Recall: Important for imbalanced datasets.

F1-score: Balances precision and recall for better assessment.

Confusion Matrix: Analyzes correct vs. incorrect classifications per digit.

5. Model Comparison

Model	Accuracy (%)	Precision	Recall	F1-Score
-----	-----	-----	-----	-----
Random Forest	94.5	High	Medium	Balanced
KNN	92.1	Medium	High	Balanced
CNN	99.2	Very High	Very High	Excellent
Transfer Learning	99.5	Very High	Very High	Excellent

6. Key Takeaways

CNN & Transfer Learning outperform traditional ML models, as they capture spatial dependencies in images.

Random Forest & KNN can work well with engineered features, but CNNs require minimal feature engineering.

Combining CNN with data augmentation enhances generalization, making it robust for real-world handwritten digit recognition

9. Visualization of Results & Model Insights

10. Tools and Technologies Used

11. Team Members and Contributions